

Adding Fast GPU Derived Datatype Handling to Existing MPIs

University of New Mexico CS Colloquium, May 5, 2021

Carl Pearson (Postdoc, Sandia National Labs)

Kun Wu, Wen-Mei Hwu, I-Hsin Chung, Jinjun Xiong



I ILLINOIS

Electrical & Computer Engineering

GRAINGER COLLEGE OF ENGINEERING

Carl Pearson



- Postdoctoral Appointee, Center for Computing Research, Sandia National Labs
 - Multi-GPU Communication
 - Sparse matrix operations
- Ph.D from University of Illinois Electrical and Computer Engineering



cwpearson



cwpears@sandia.gov



carlpearson.net

Outline

- Stencil code
 - Domain Decomposition
- Stencil code and supercomputing
 - Processes, MPI, and GPUs
- Non-contiguous data
 - Where it comes from and why it matters
- TEMPI's approach to derived type handling
 - Translation
 - Canonicalization
 - Kernel Selection
- Some Performance Results
- How TEMPI works with MPI

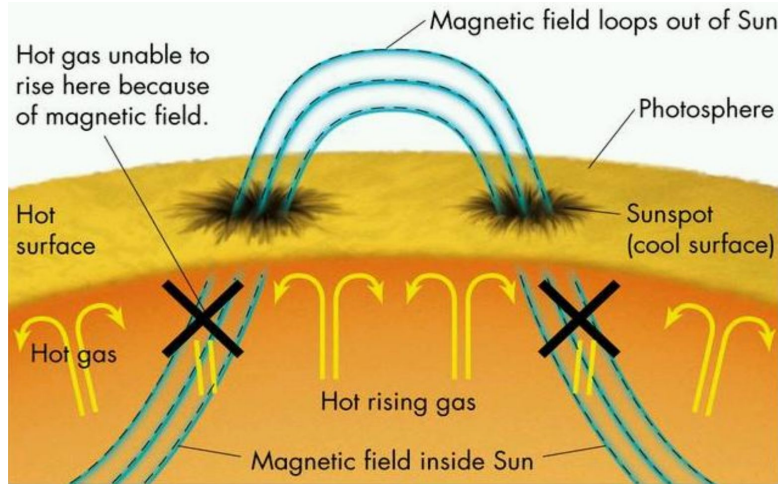
CS 375

CS 341, 442, 471, 475, 481

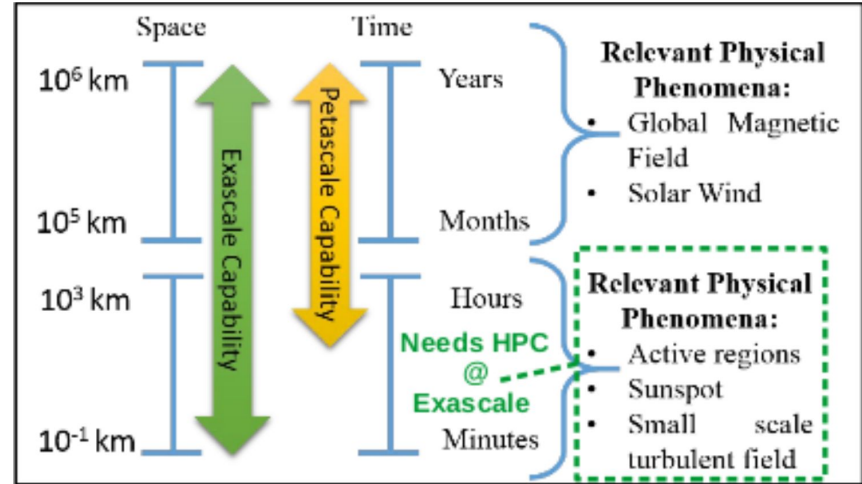
CS 151, 241, 251

Astaroth¹ 3D Stellar Simulation Code

- Aalto Univ. / Univ. of Helsinki
- 2E+11 gridpoints on 3072 GPUS (256³ / GPU)
- First direct numerical simulation of relevant high-resolution phenomena

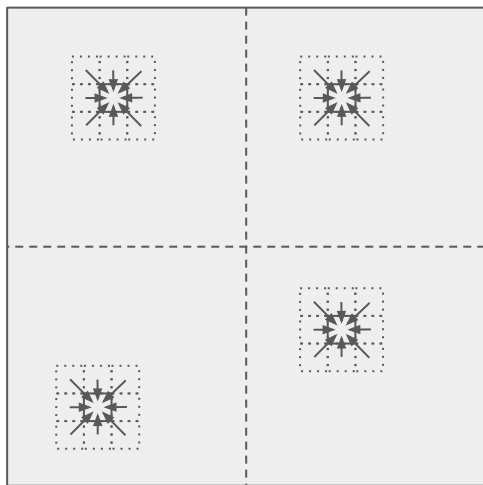


Figures from a talk by Anjum, O. 2017

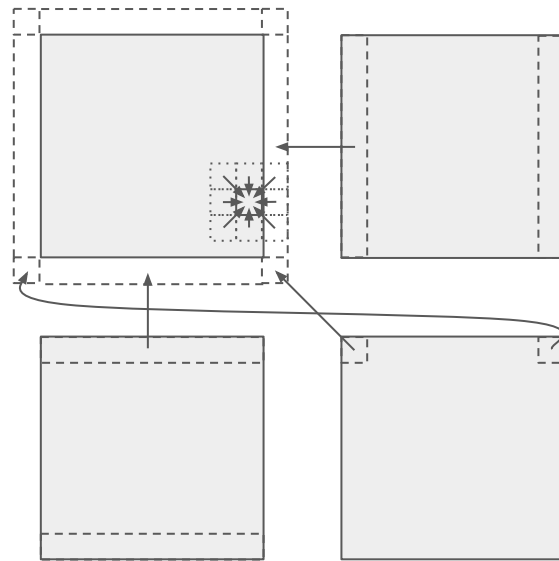


Distributed Stencil

- Inside of sun as a grid of points
 - Each point has associated physical properties
- Value at next timestep computed from neighbor's current values

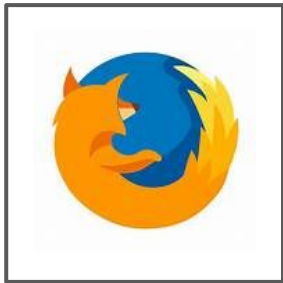


domain
decomposition of
grid



Sub-grids with boundary
communication

Parallel Computing with Processes

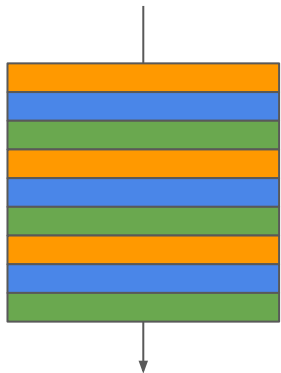


- Program code
- Program state
- Isolation
 - Each process thinks it is alone on the computer
 - Errors in one process do not affect others

Parallel Computing with Processes



- Program code
- Program state
- Isolation
 - Each process thinks it is alone on the computer
 - Errors in one process do not affect others

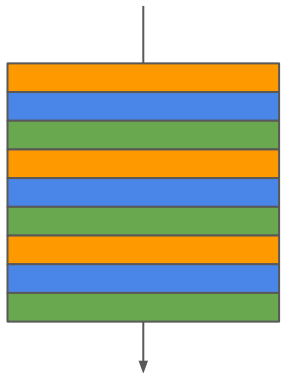


One CPU

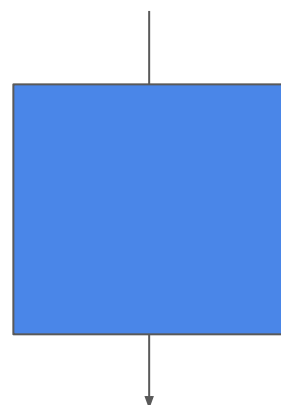
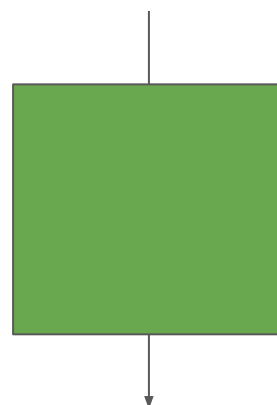
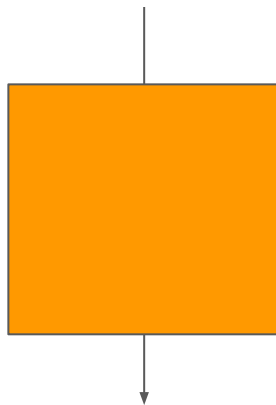
Parallel Computing with Processes



- Program code
- Program state
- Isolation
 - Each process thinks it is alone on the computer
 - Errors in one process do not affect others



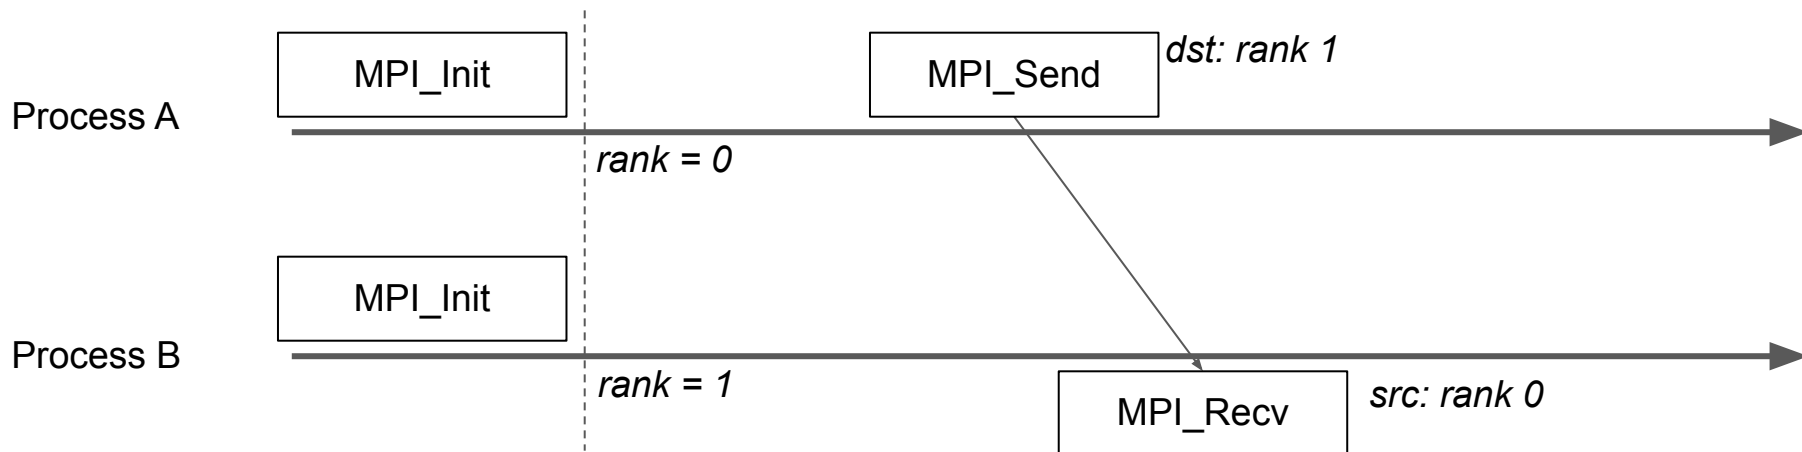
One CPU



Three CPUs (more done in the same amount of time)

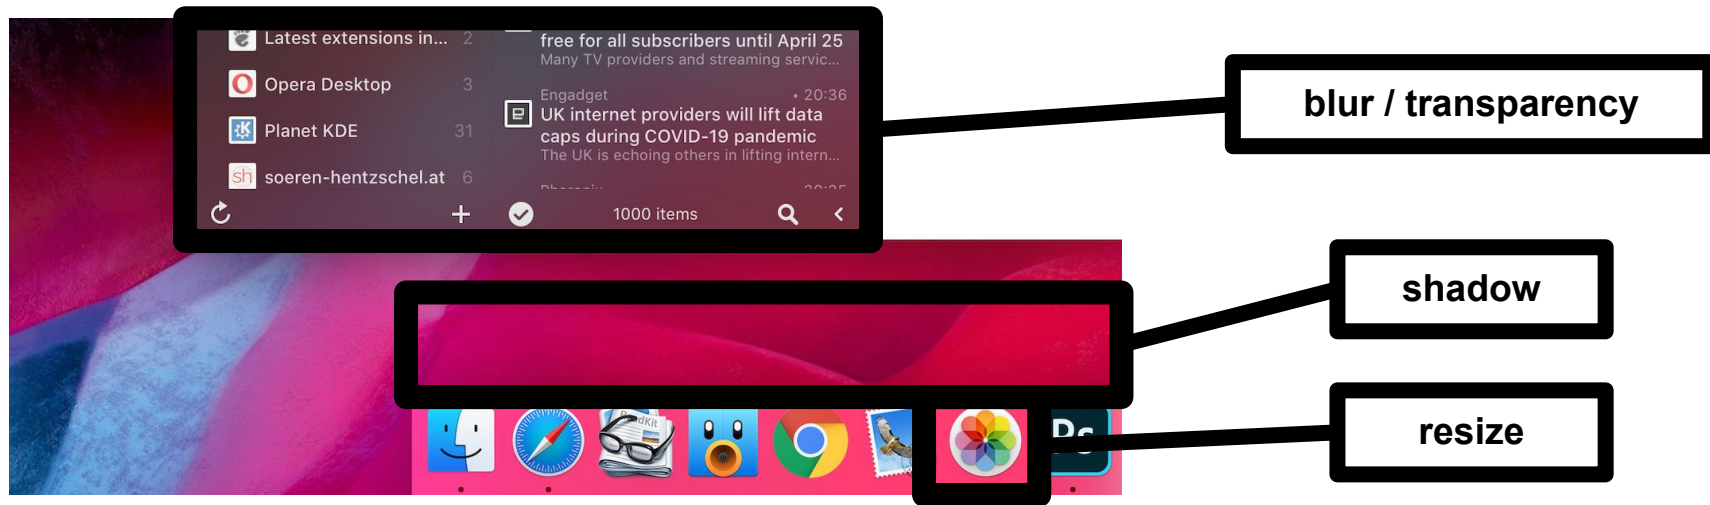
MPI - Message Passing Interface

- Selectively break isolation between processes
- Coordination: “have you given me updated boundary values?”
- Communication: “here are updated boundary values for you.”



GPU Computing (Graphics Processing Unit)

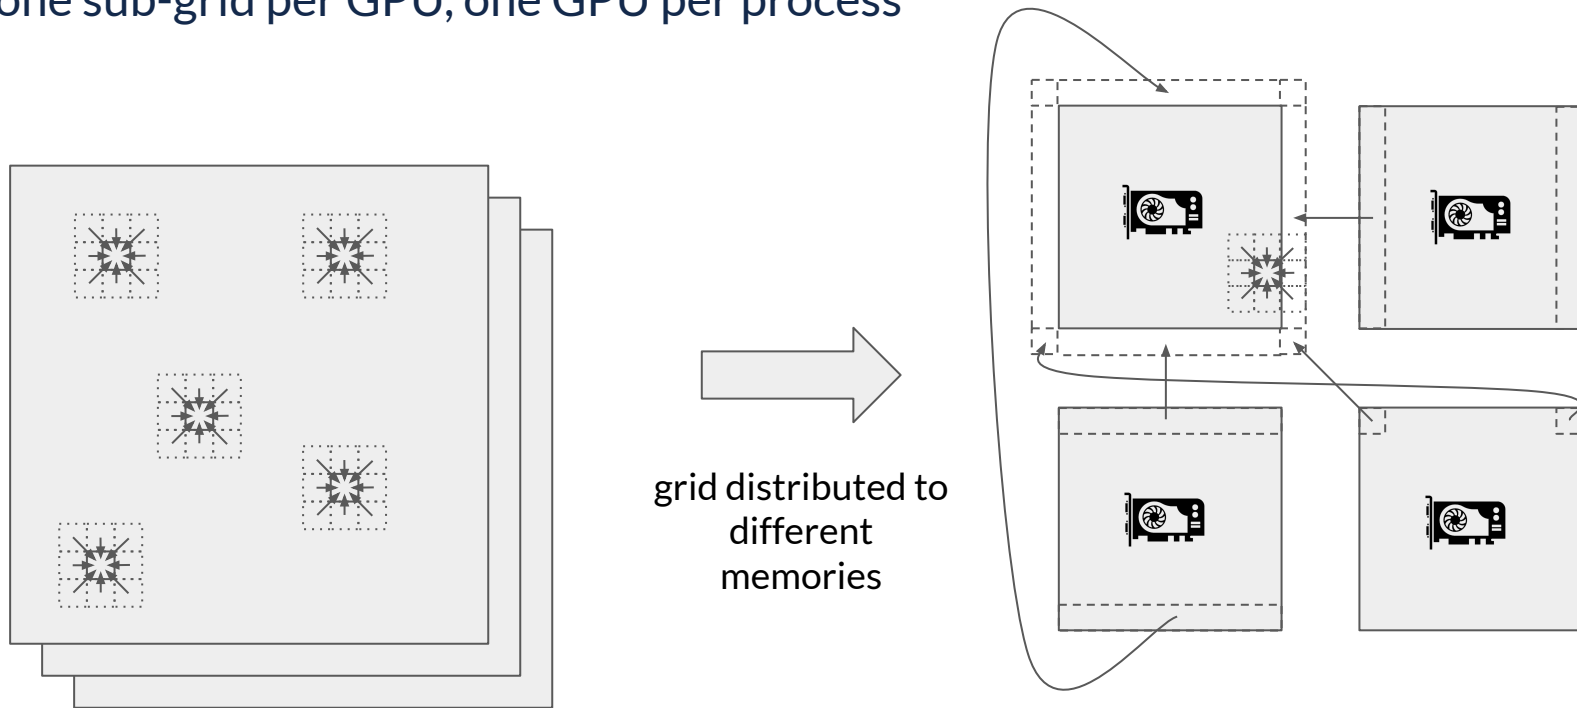
- Special hardware to render displayed images faster
 - Quickly apply same operation to many pixels



- Turns out, many scientific codes have similar properties
- Blur: create new pixel value based on neighbors
- Stencil: create new gridpoint value based on neighbors

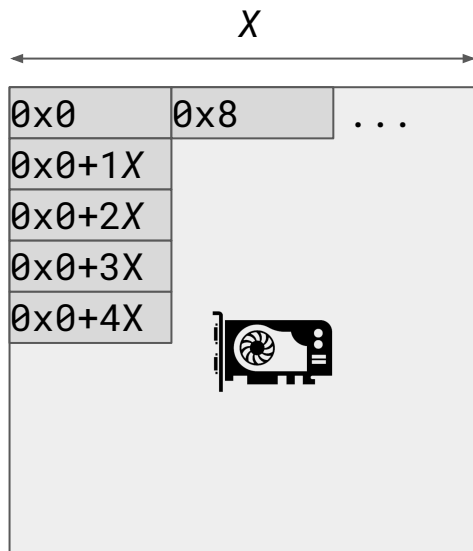
Distributed Stencil

- one sub-grid per GPU, one GPU per process



Contiguous & Non-contiguous Data

- “row-major” storage



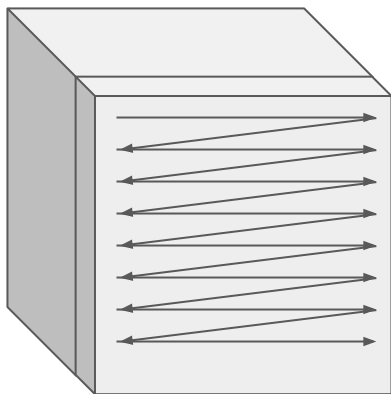
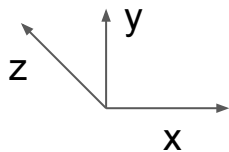
“top edge” - contiguous



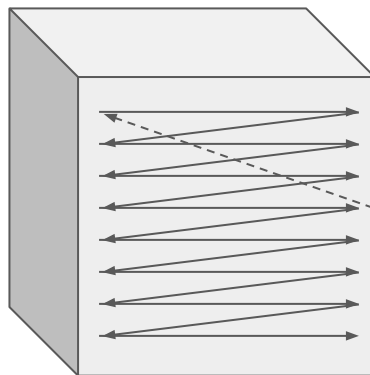
“left edge” - non-contiguous



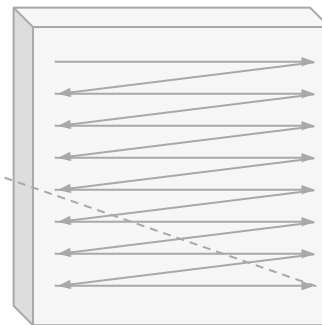
“Row-major” Storage in 3D



$z = 0$ plane



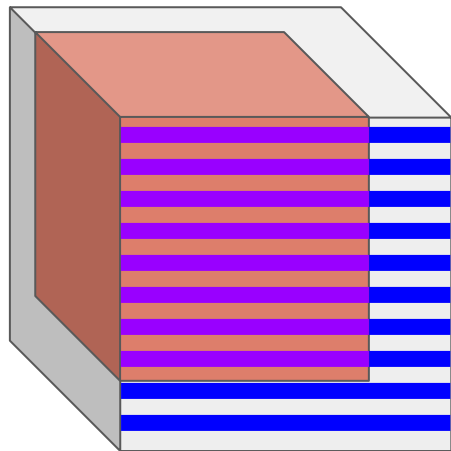
$z = 1$ plane



$z = 0$ plane

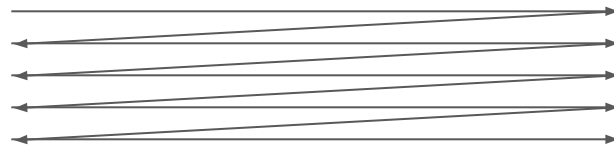
3D Non-contiguous Example

3D Region



$(z = 0)$

(storage order)

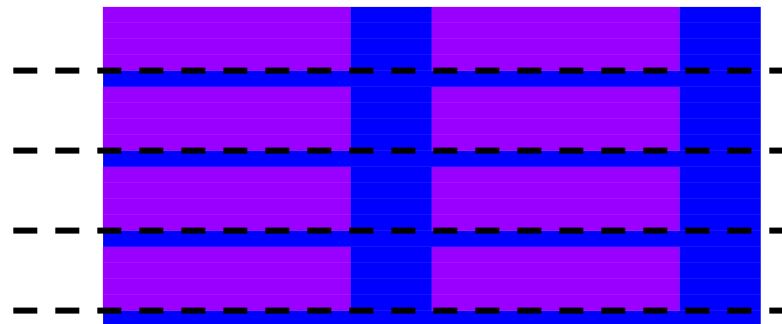


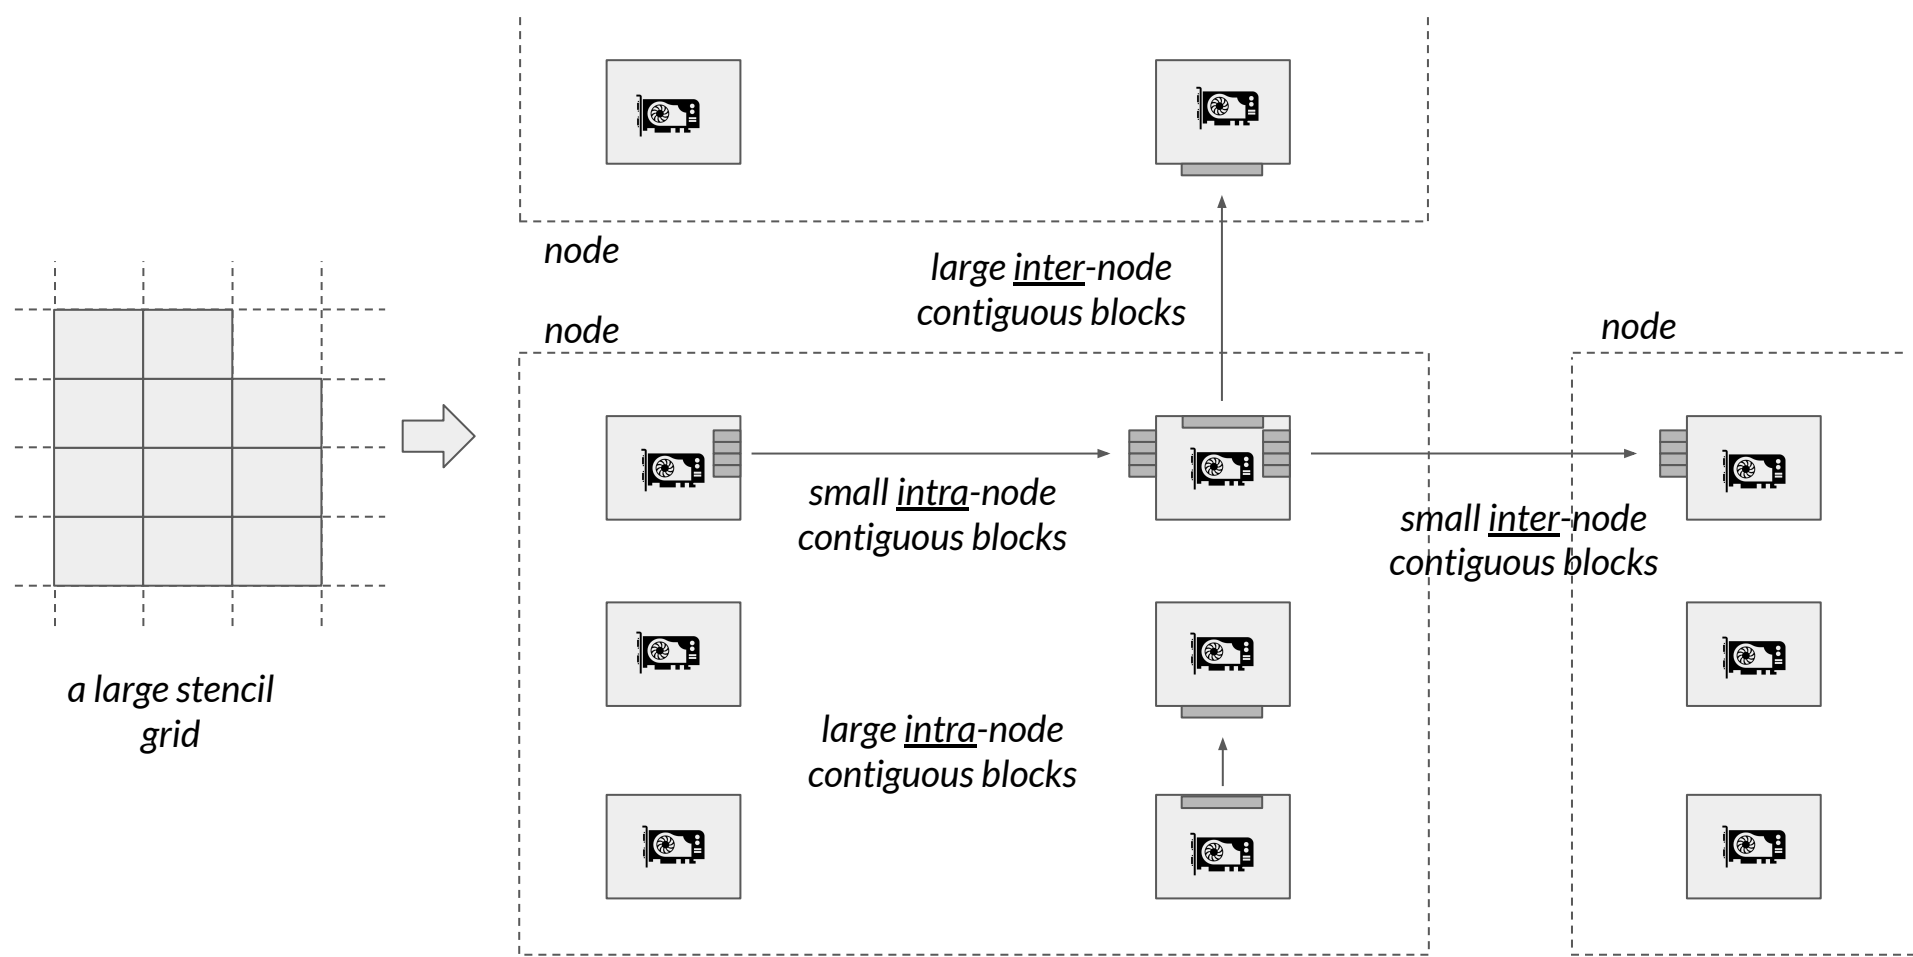
$z=0$

$z=1$

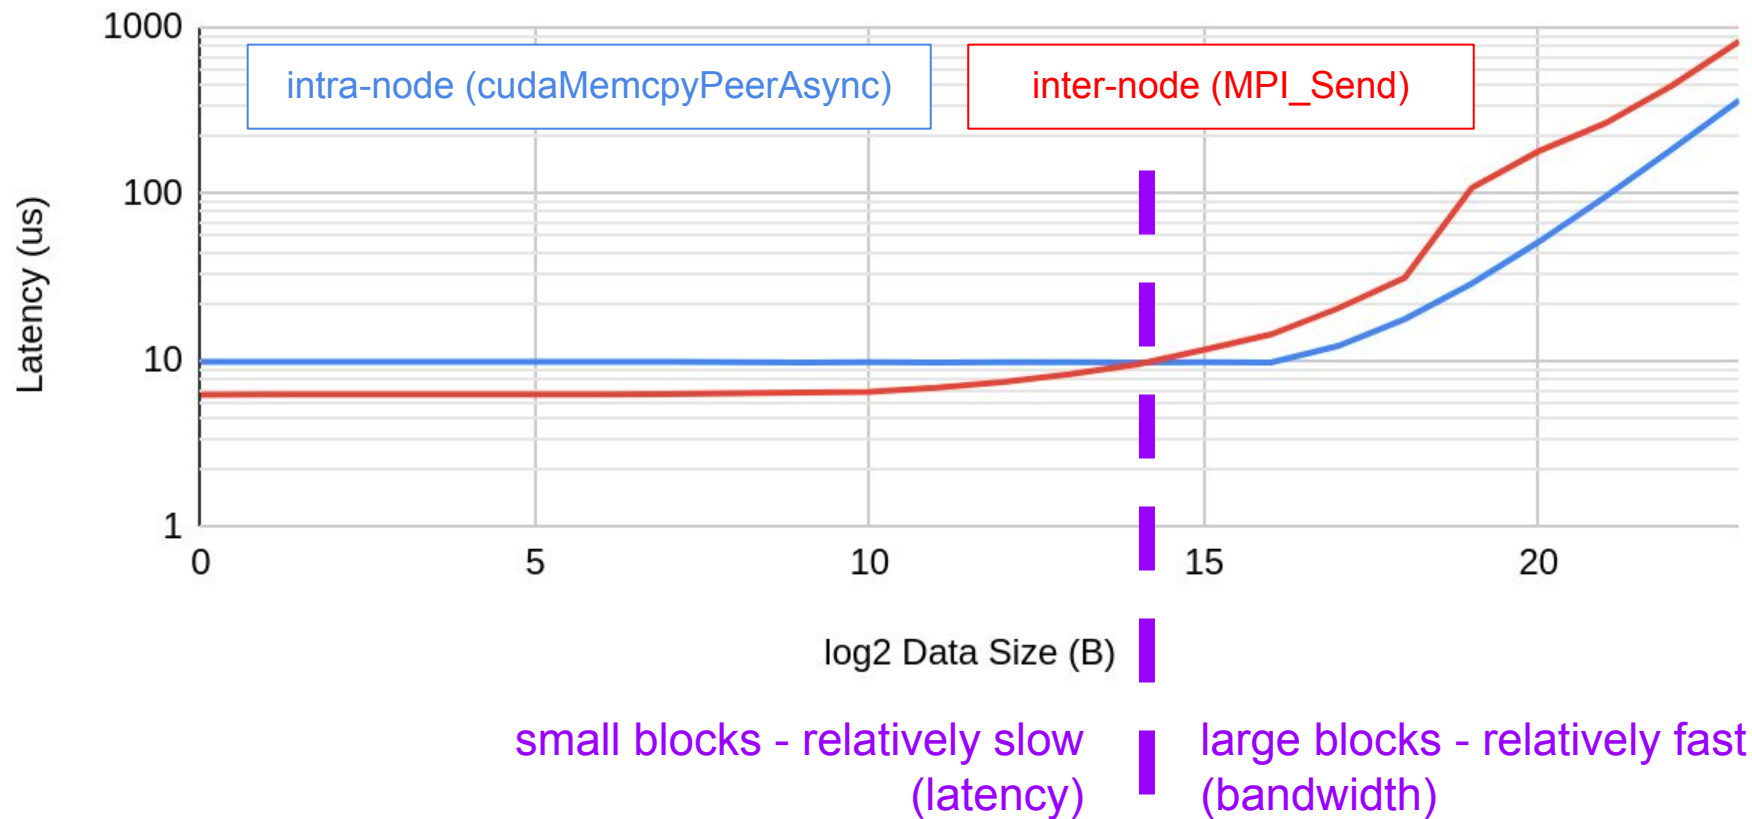
$z=2$

...

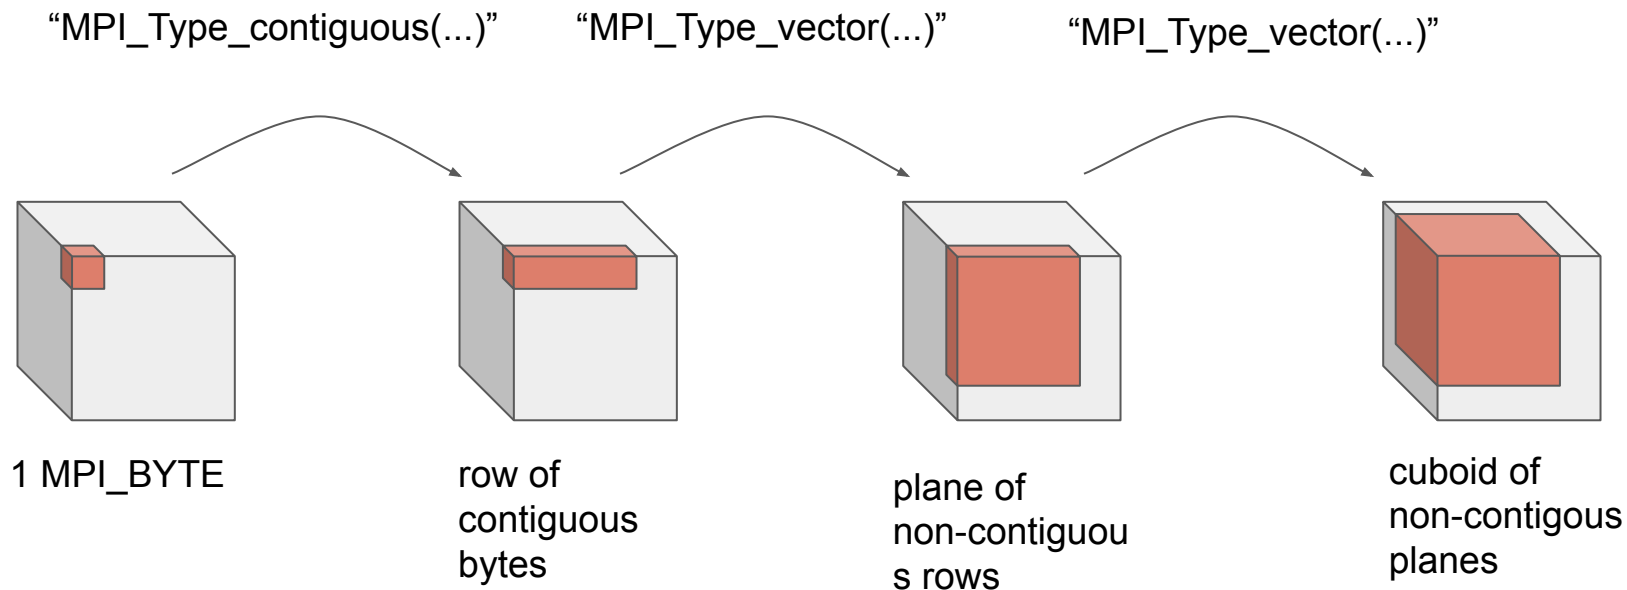




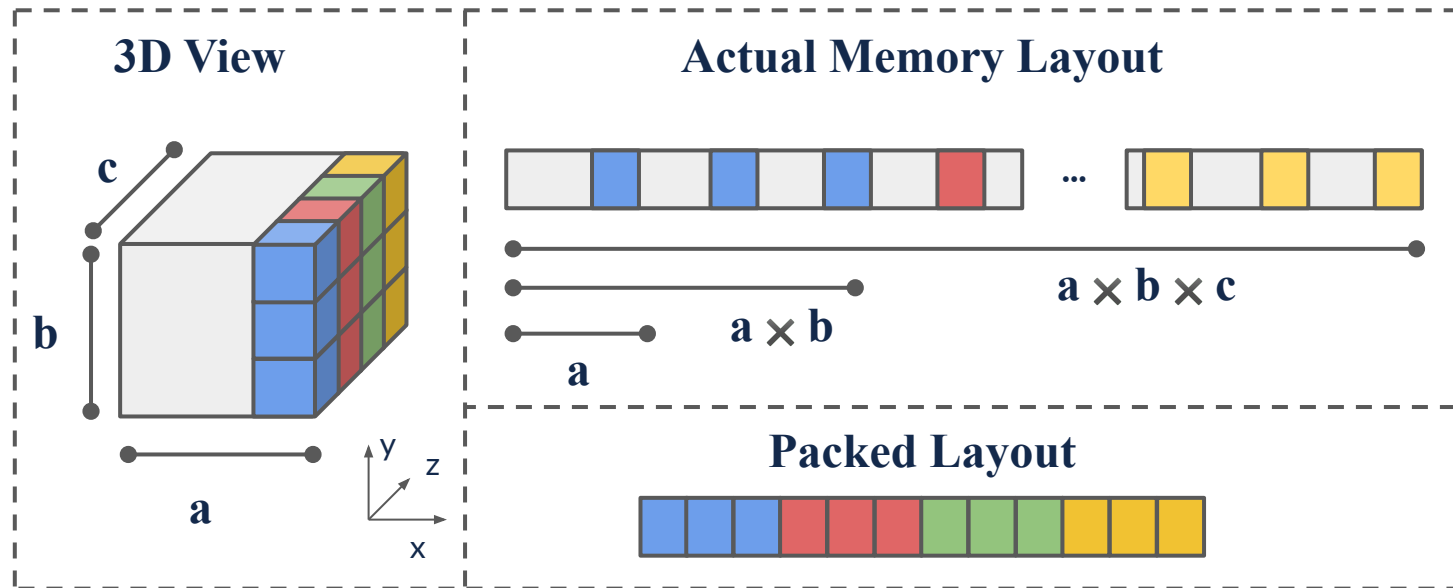
Latency vs Contiguous Size



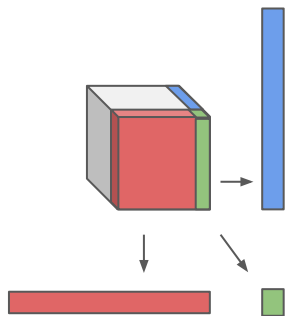
MPI Derived Datatypes



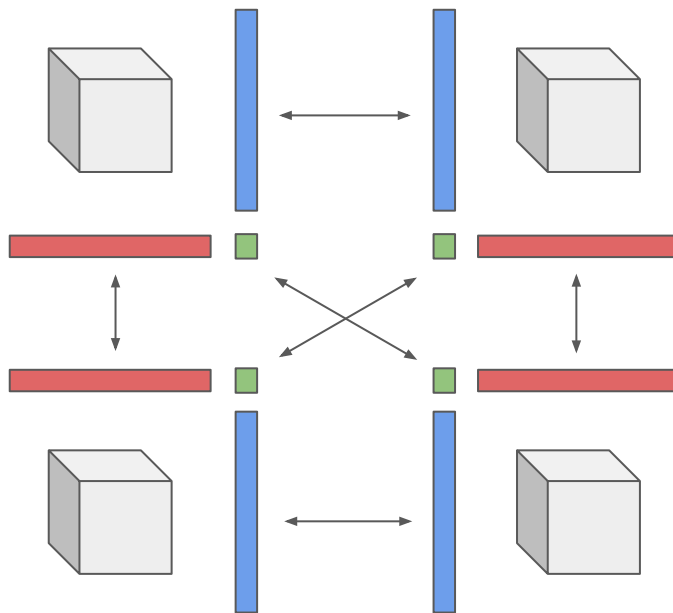
Many Small Blocks into Few Large Blocks



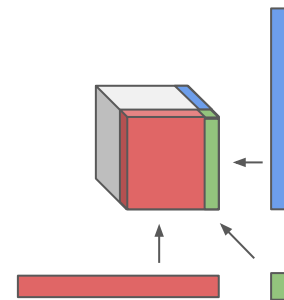
Pack / Alltoallv / Unpack



MPI_Packs

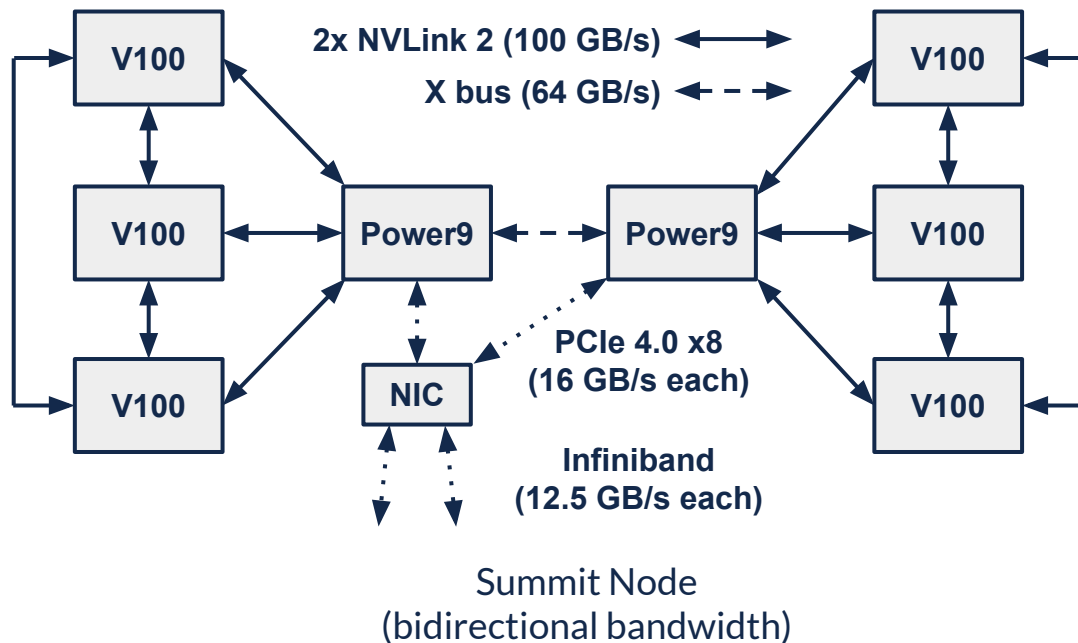


MPI_Neighbor_alltoallv



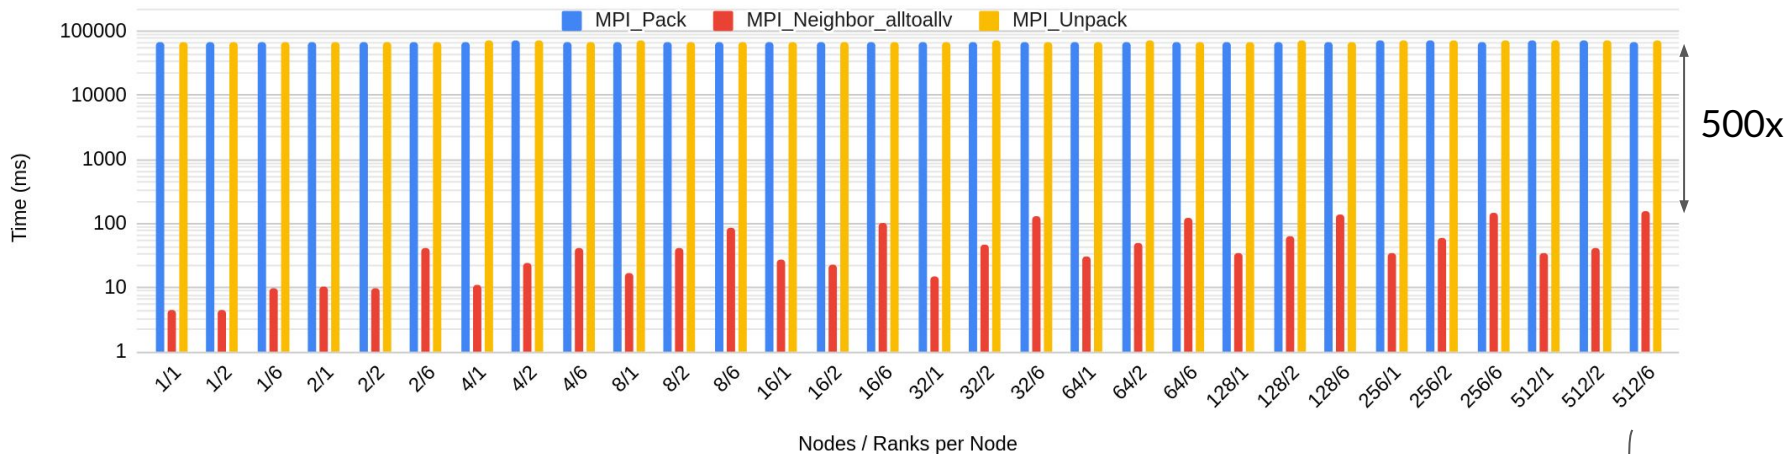
MPI_Unpacks

OLCF Summit



The Problem (1/2)

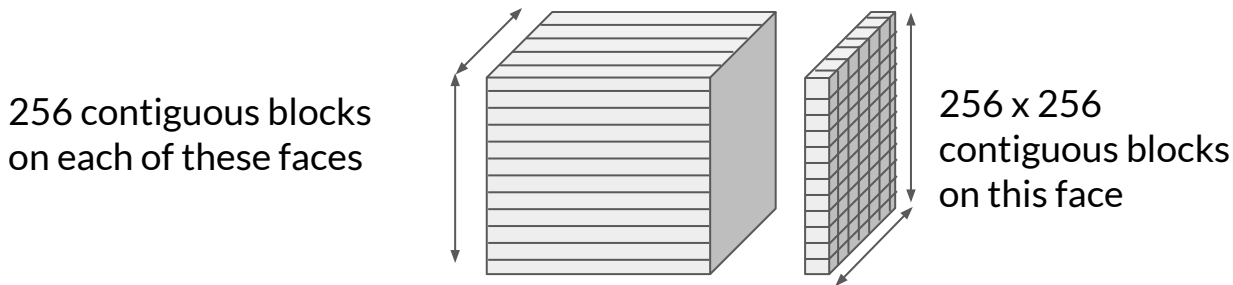
- Halo exchange with MPI derived types



- 73.7 MiB/rank
 - MPI_Neighbor_alltoallv = ~500 MB/s/rank
 - MPI_Pack / MPI_Unpack = ~1 MB/s

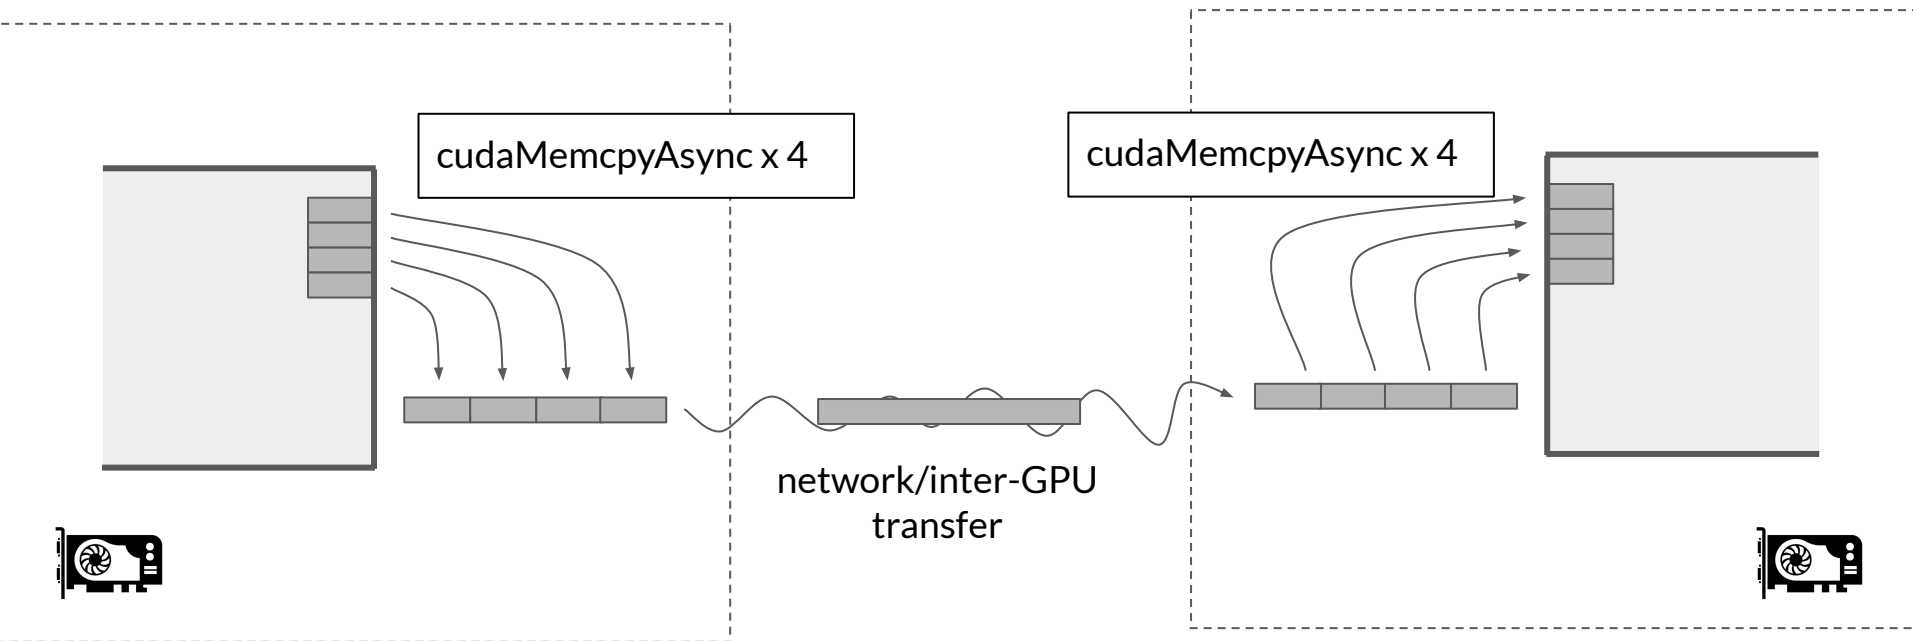
The Problem (2/2)

- Most of the “non-contiguousness” is in one dimension
- 3,145,728 contiguous blocks (~20us per block)
- one `cudaMemcpyAsync` per block



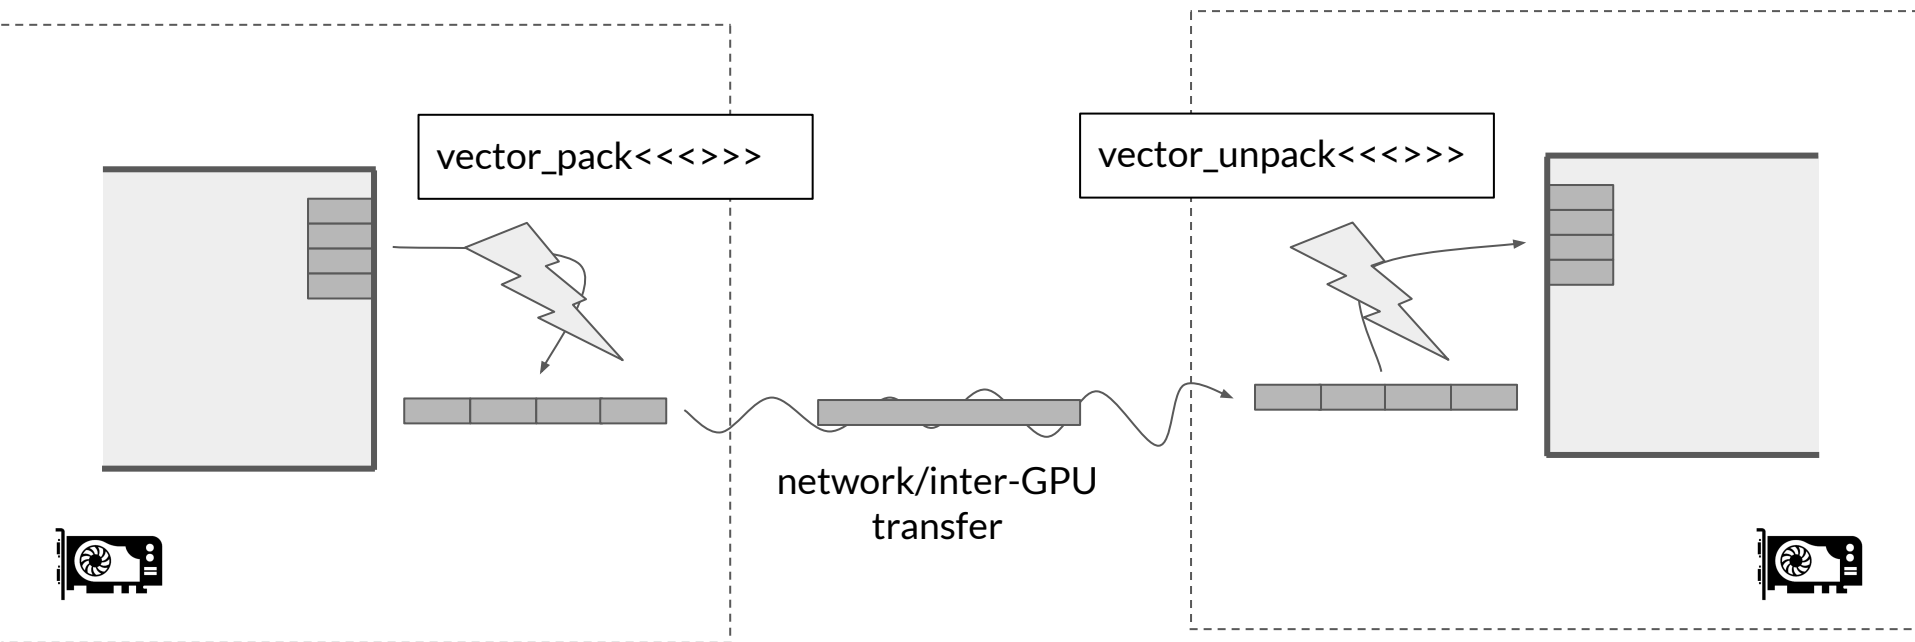
(recall that halo space breaks up otherwise contiguous directions)

MPI_Send (SpectrumMPI)

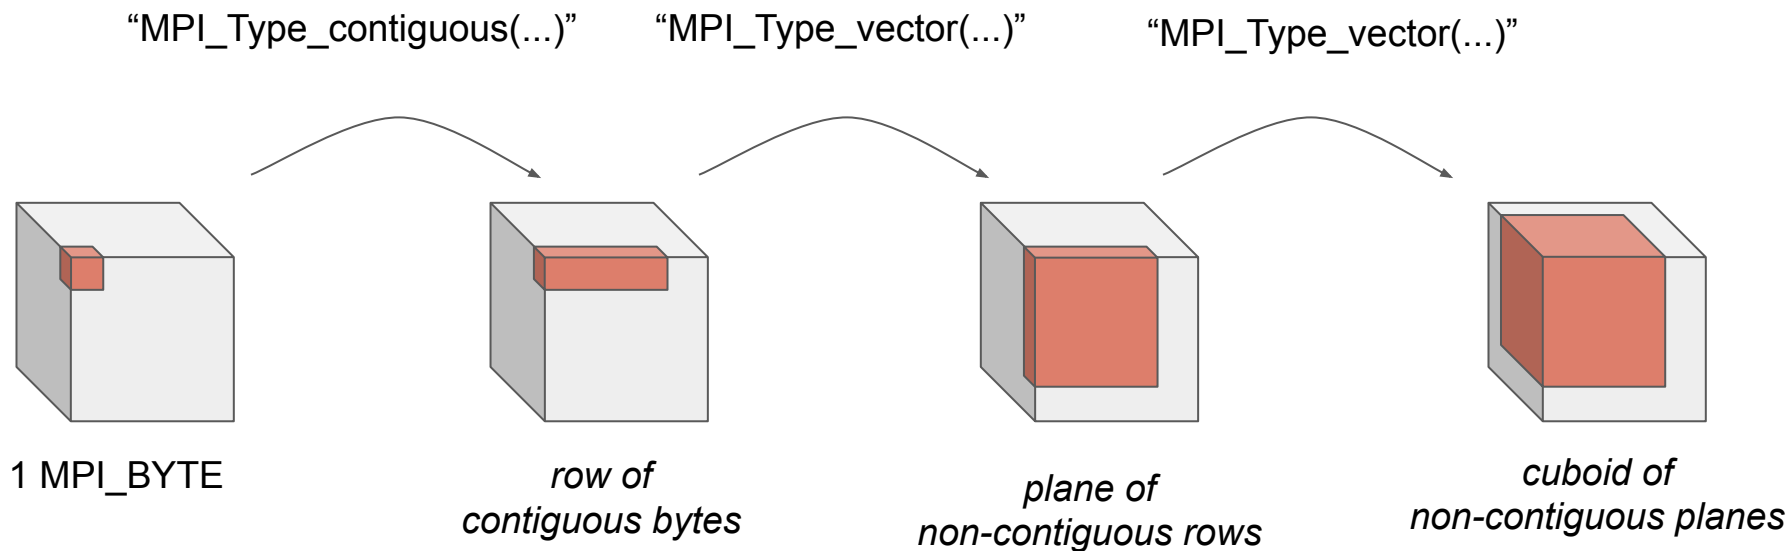


A Solution

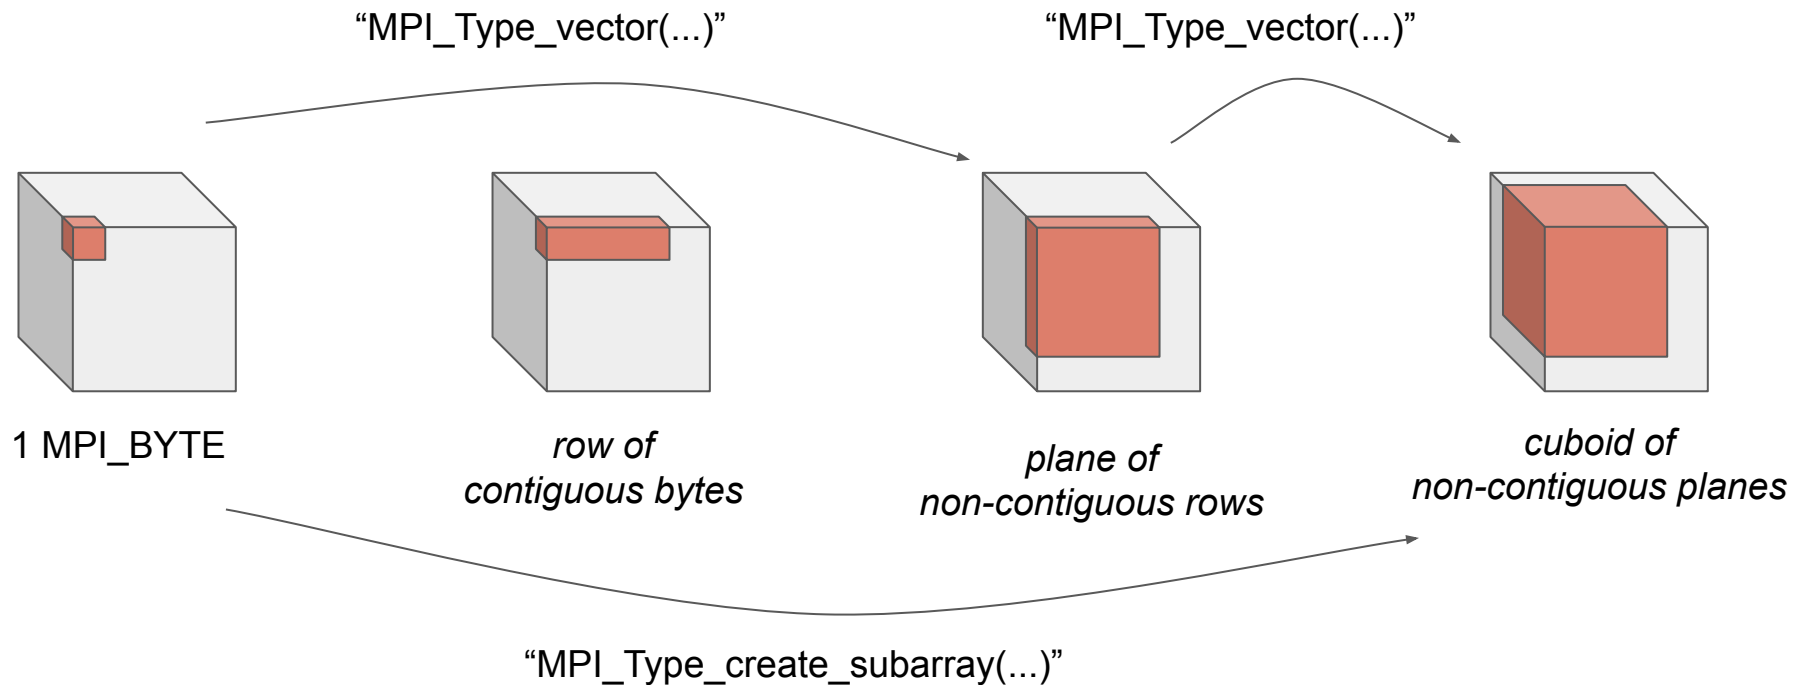
- GPU Kernels to pack non-contiguous data



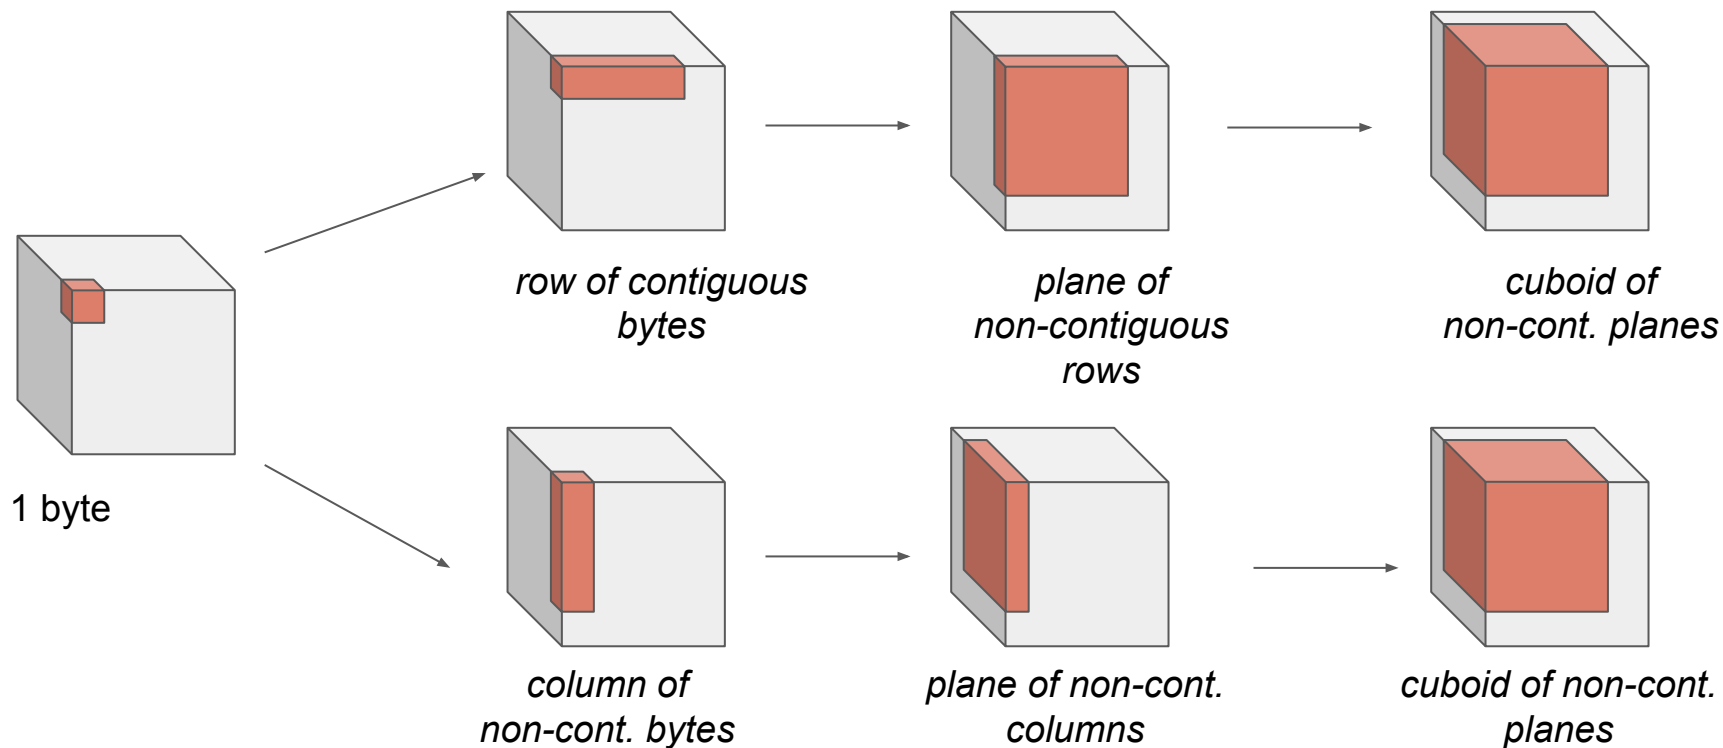
MPI Derived Datatypes: Revisited



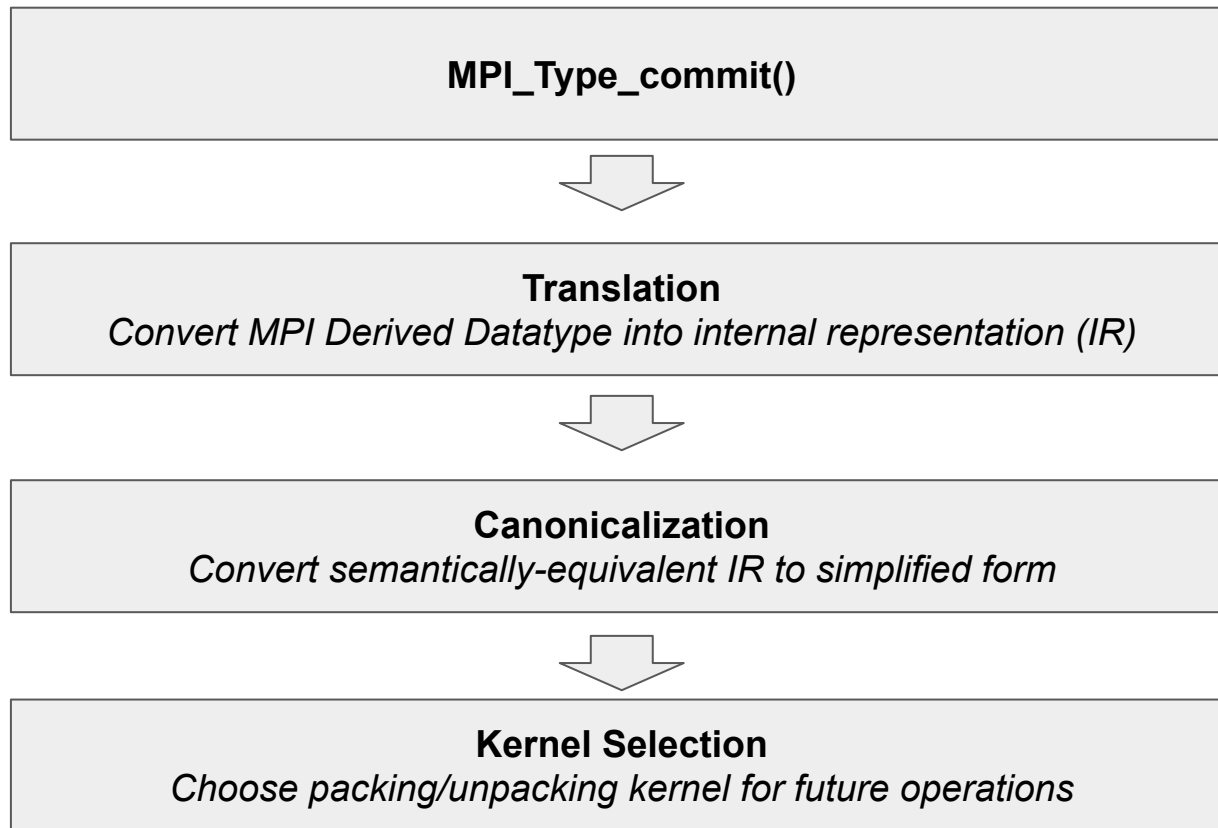
MPI Derived Datatypes: Revisited



MPI Derived Datatypes: Revisited



TEMPI Datatype Handling

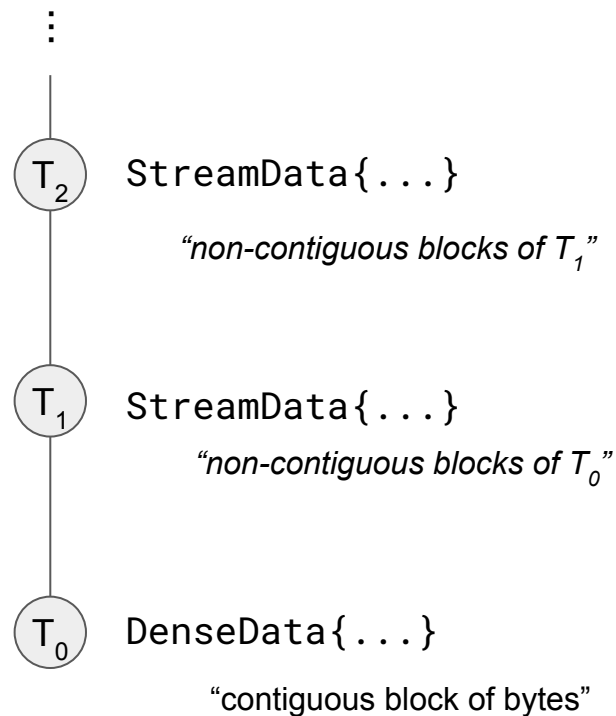


IR - “Internal Representation”

```
StreamData {  
    integer offset; // offset (B) of the first element  
    integer stride; // pitch (B) between element  
    integer count;  // number of elements  
}
```

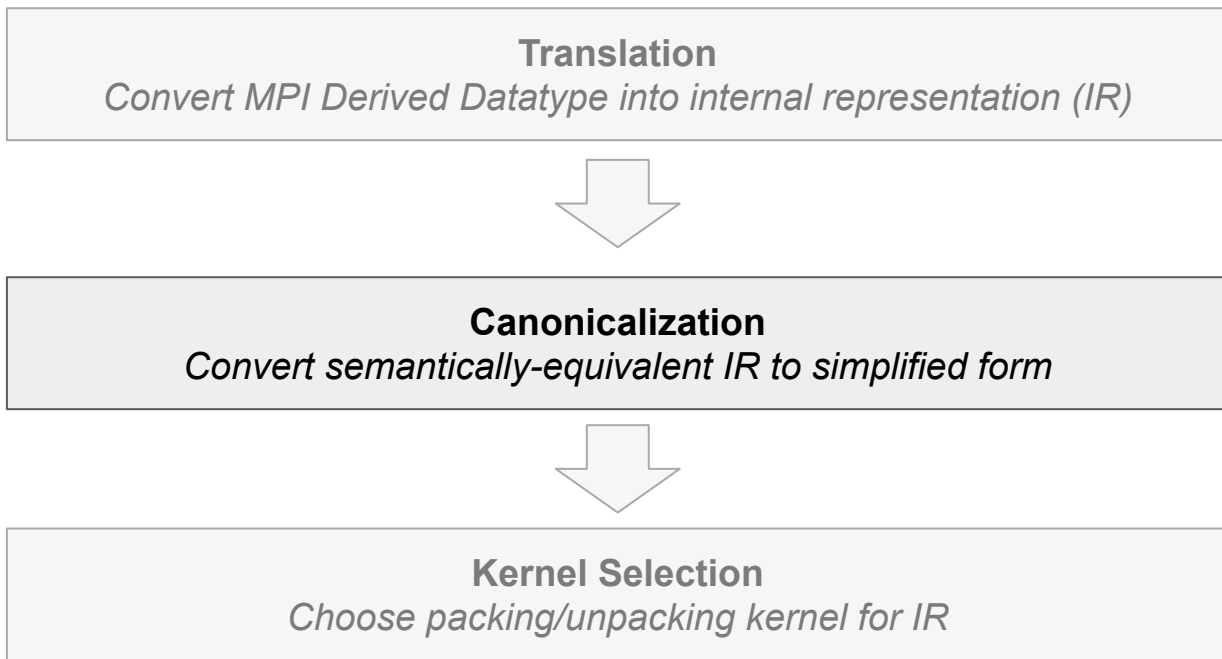
```
DenseData {  
    integer offset; // offset (B) of the first byte  
    integer extent; // number of bytes  
}
```

Hierarchy of StreamData, rooted at DenseData

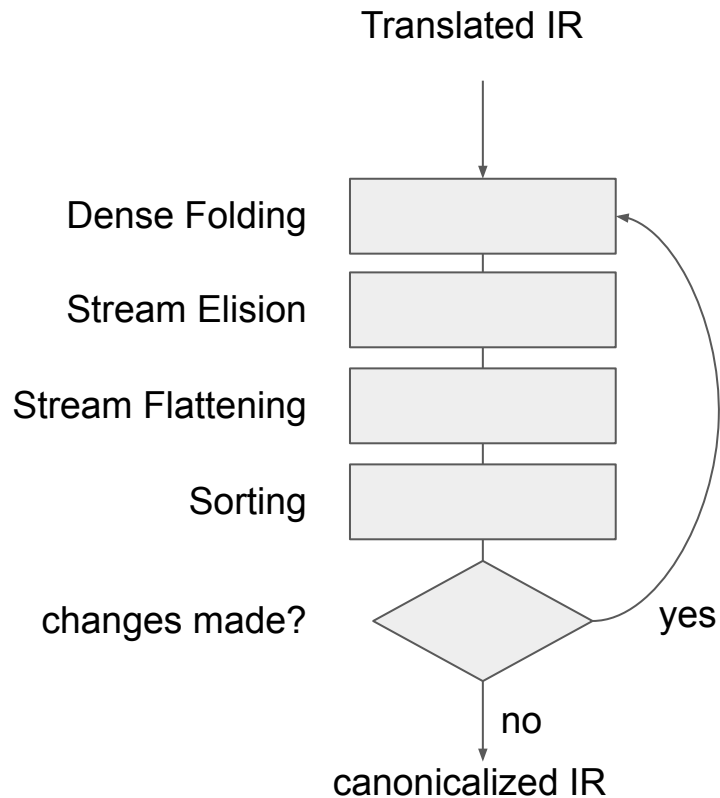


Examples

TEMPI Datatype Handling

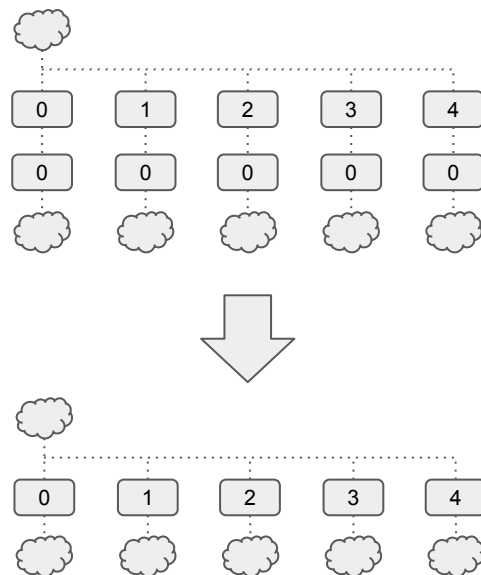
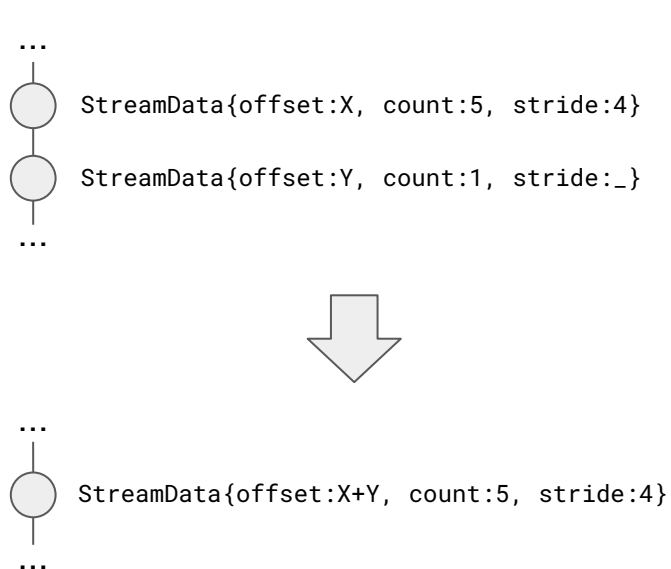


Canonicalization



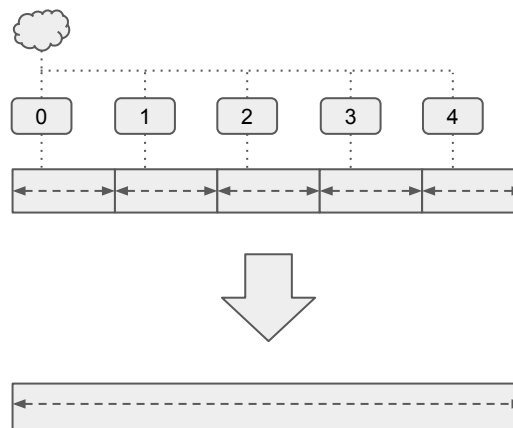
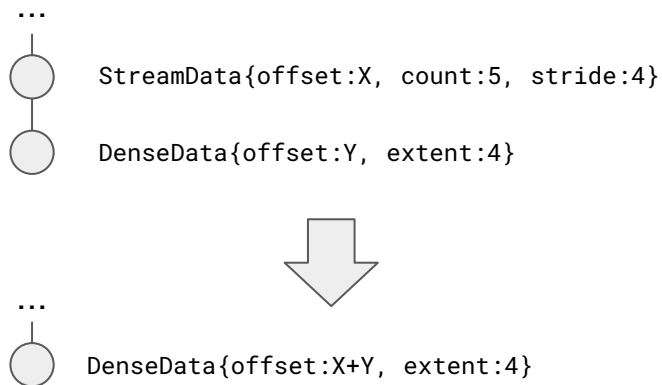
Canonicalization: Stream Elision

An MPI vector will commonly have a block of 1 child element



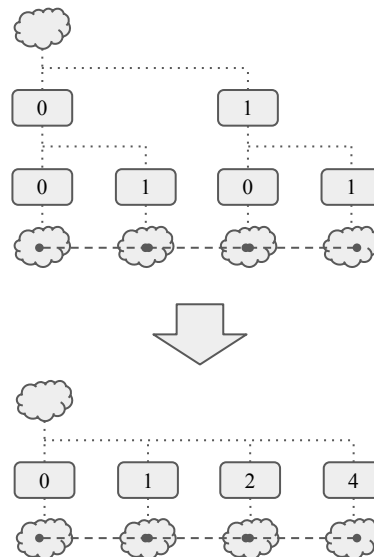
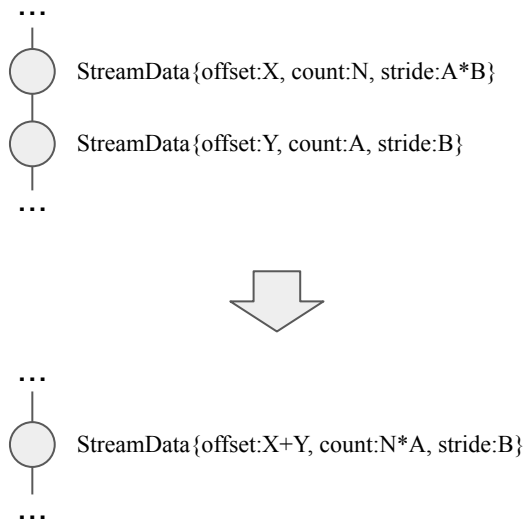
Canonicalization: Dense Folding

When a stream is actually multiple dense elements
A parent type of an MPI named type

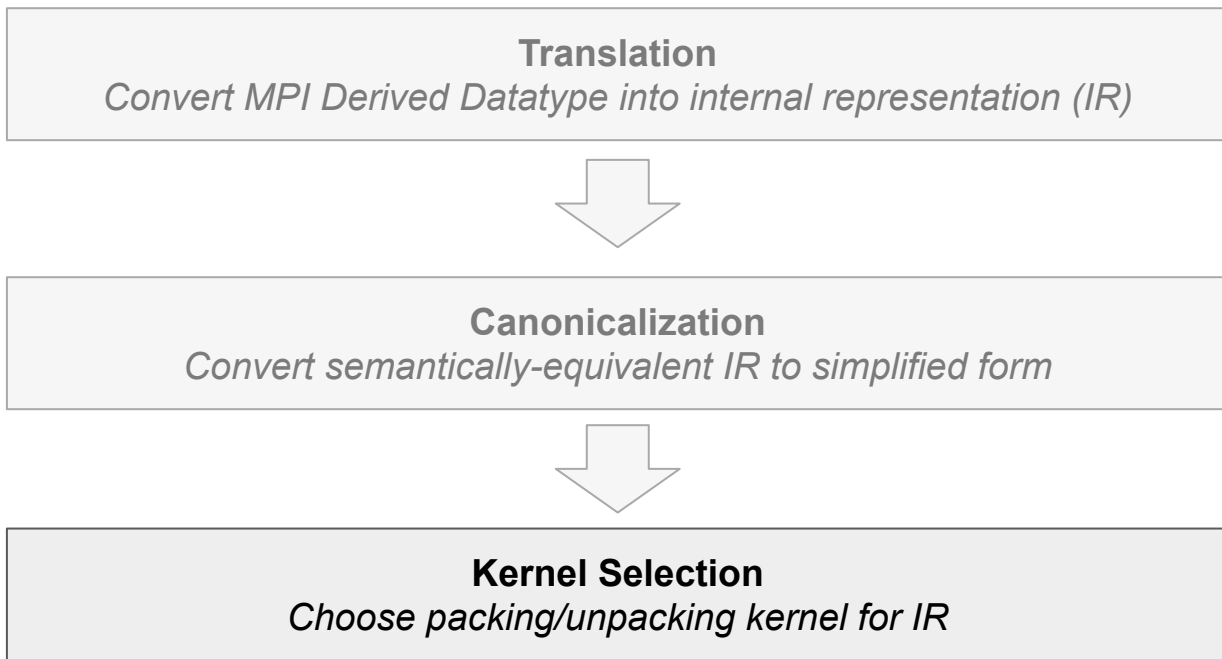


Canonicalization: Stream Flattening

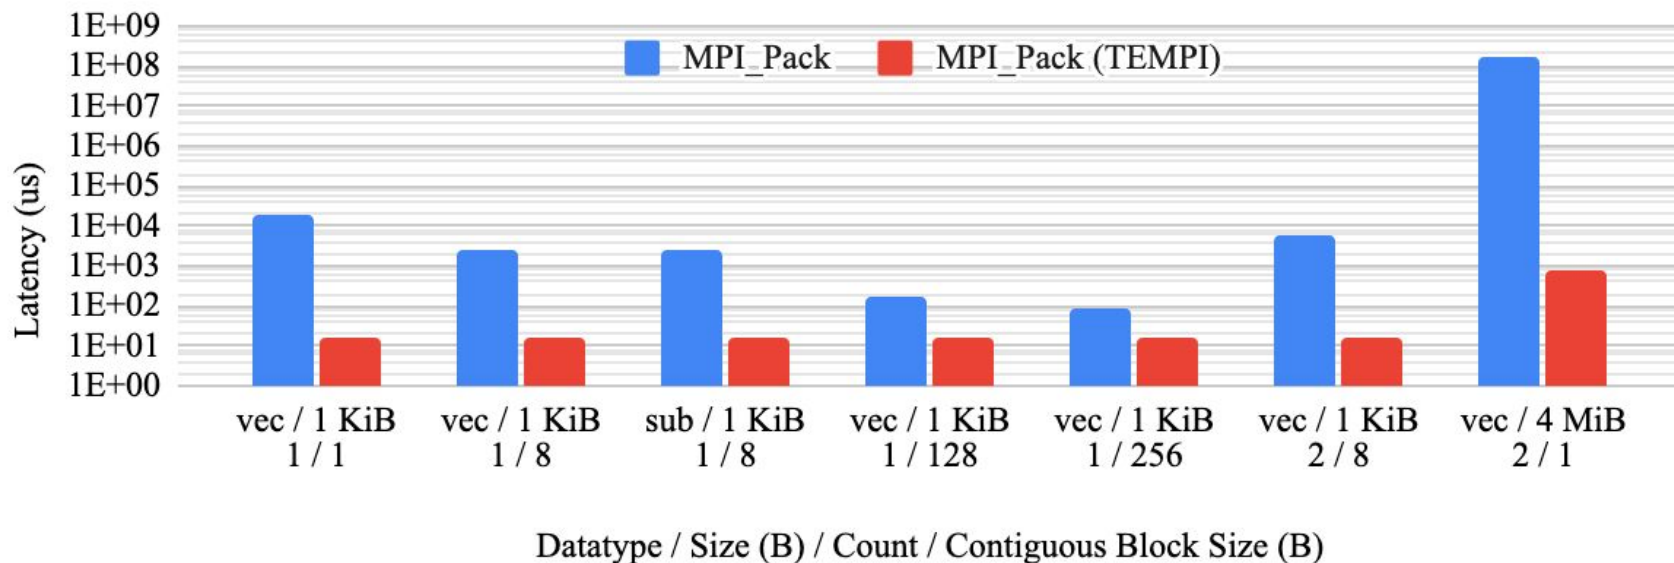
When two streams of three elements is one stream of six elements



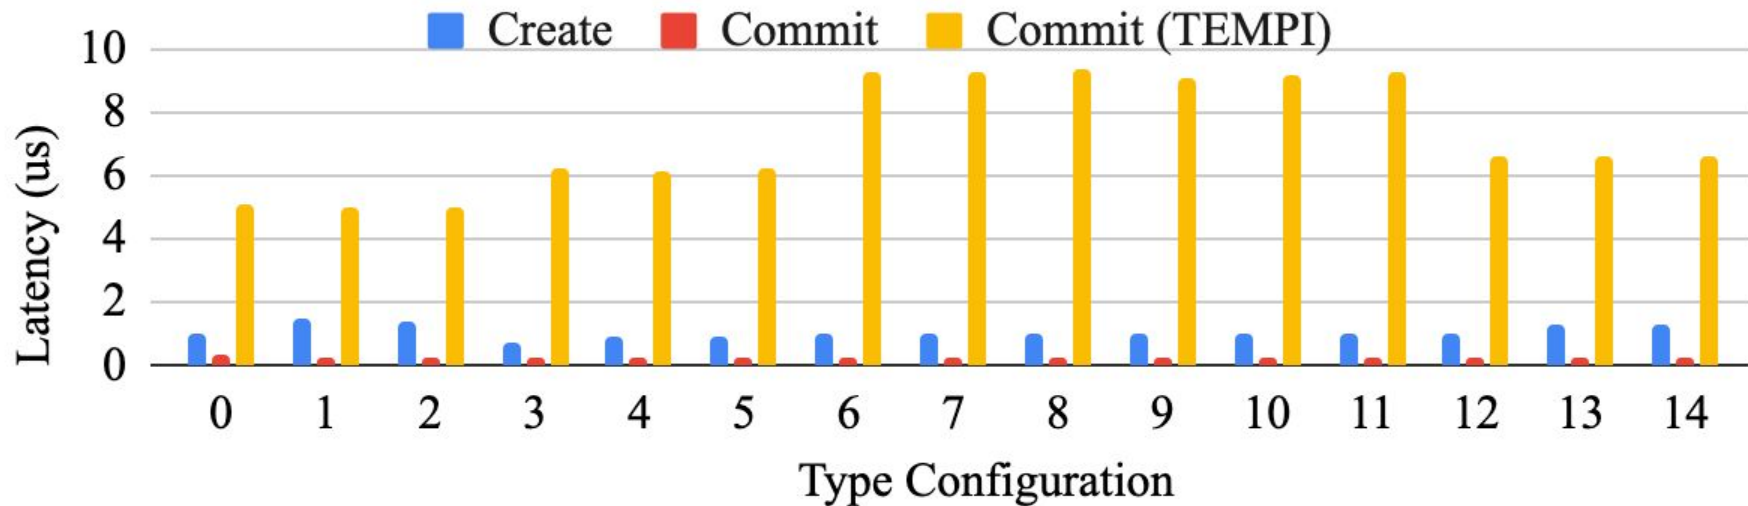
TEMPI Datatype Handling



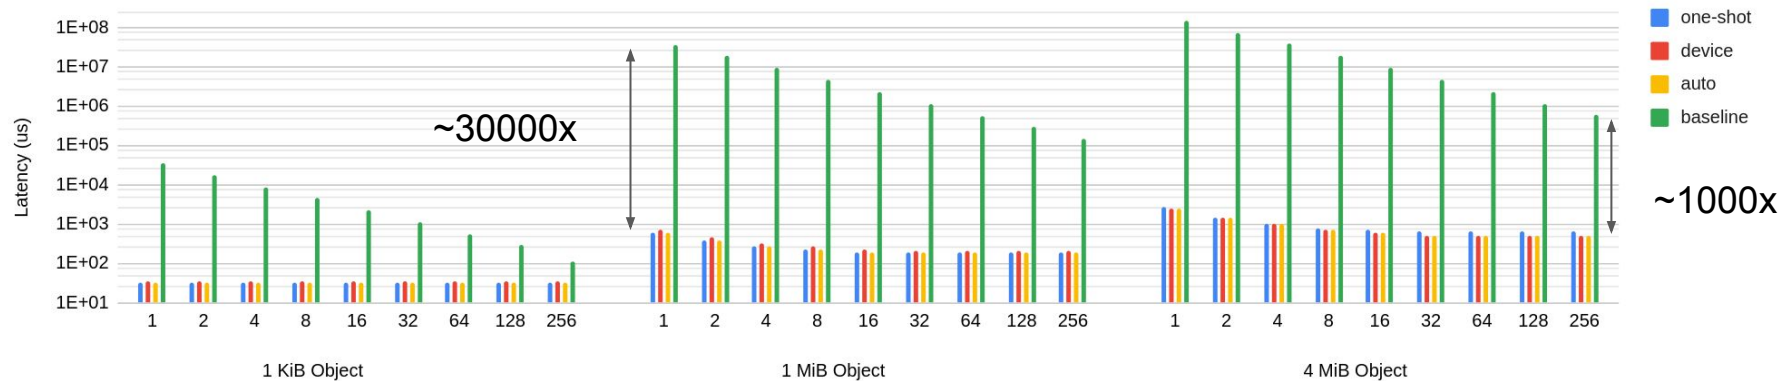
MPI_Pack Results



MPI_Type_commit Time

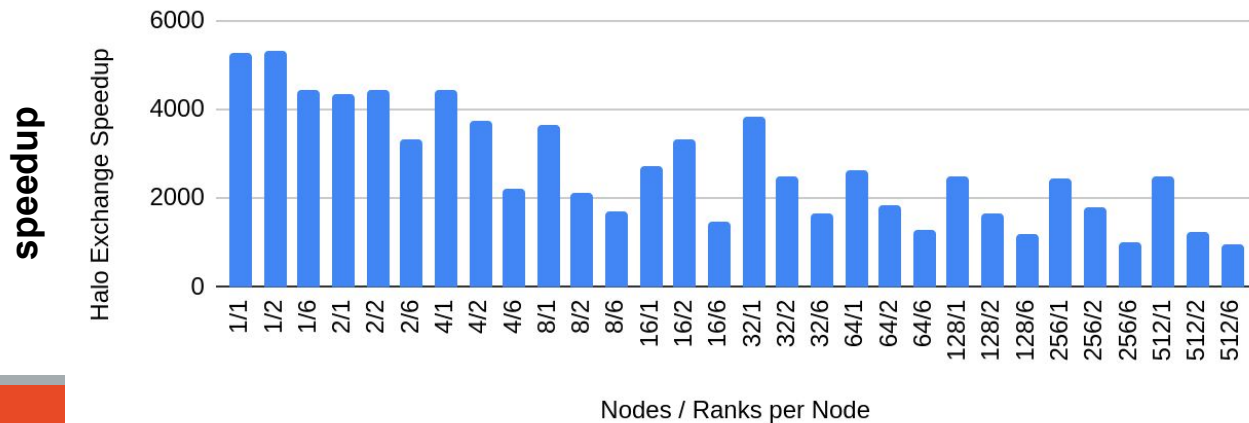
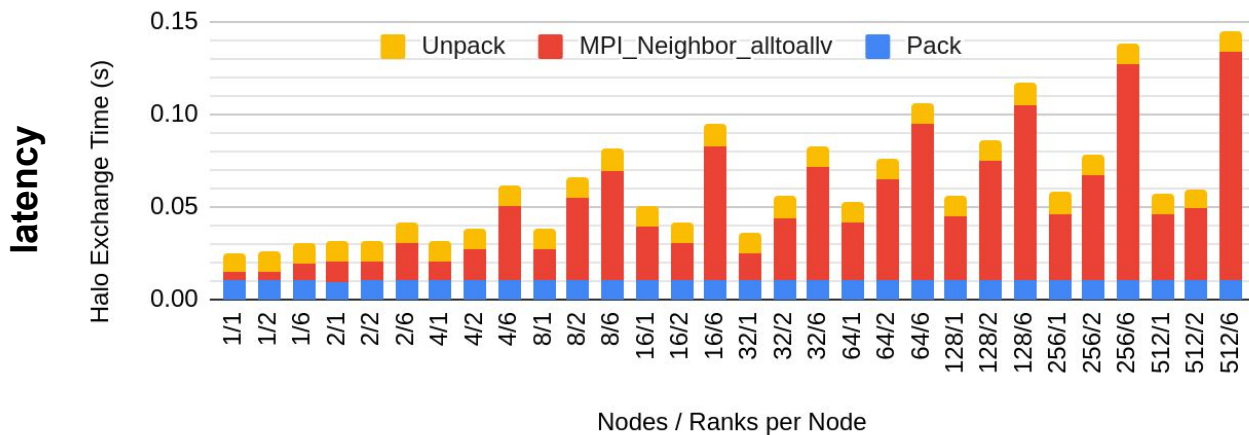


MPI_Send / MPI_Recv



MPI_Send/Recv Latency for 2D objects with different block sizes

Halo Exchange

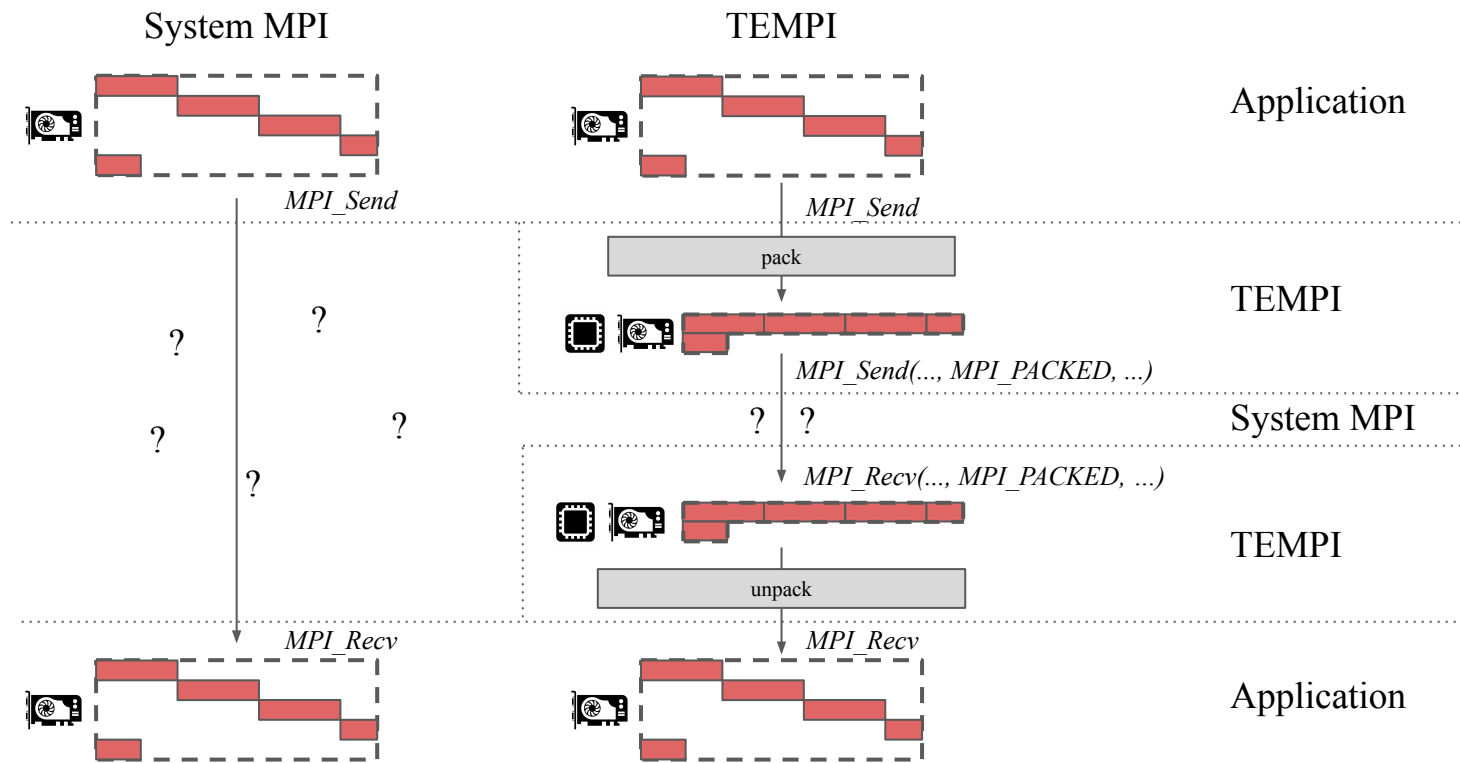


A Practical Challenge

- Large-scale systems are tightly controlled
 - Can't just make whatever changes you want
 - Usually one MPI (or maybe two) are deployed on the system
 - Rarely bugfixed (if ever)
 - Even more rarely are new features added
 - Difficult or impossible to make experimental modifications
-
- MPI has a well-defined standard
 - Take advantage of this + how OS loads libraries to inject modifications

TEMPI

- “Temporary MPI” / “Topology Experiments for MPI” / plural of tempo (speed)
- MPI interposer



app.c

```
#include <mpi.h>
```

1

```
int main(int argc, char **argv) {  
    MPI_Init(&argc, &argv);  
}
```



```
gcc app.c -o app \  
-I /mpi/include \  
-L /mpi/lib \  
-l mpi
```

2



app

```
...  
callq 7260  
<MPI_Init@plt>  
...
```

4

8

system MPI

mpi.h

```
int MPI_Init(...);
```

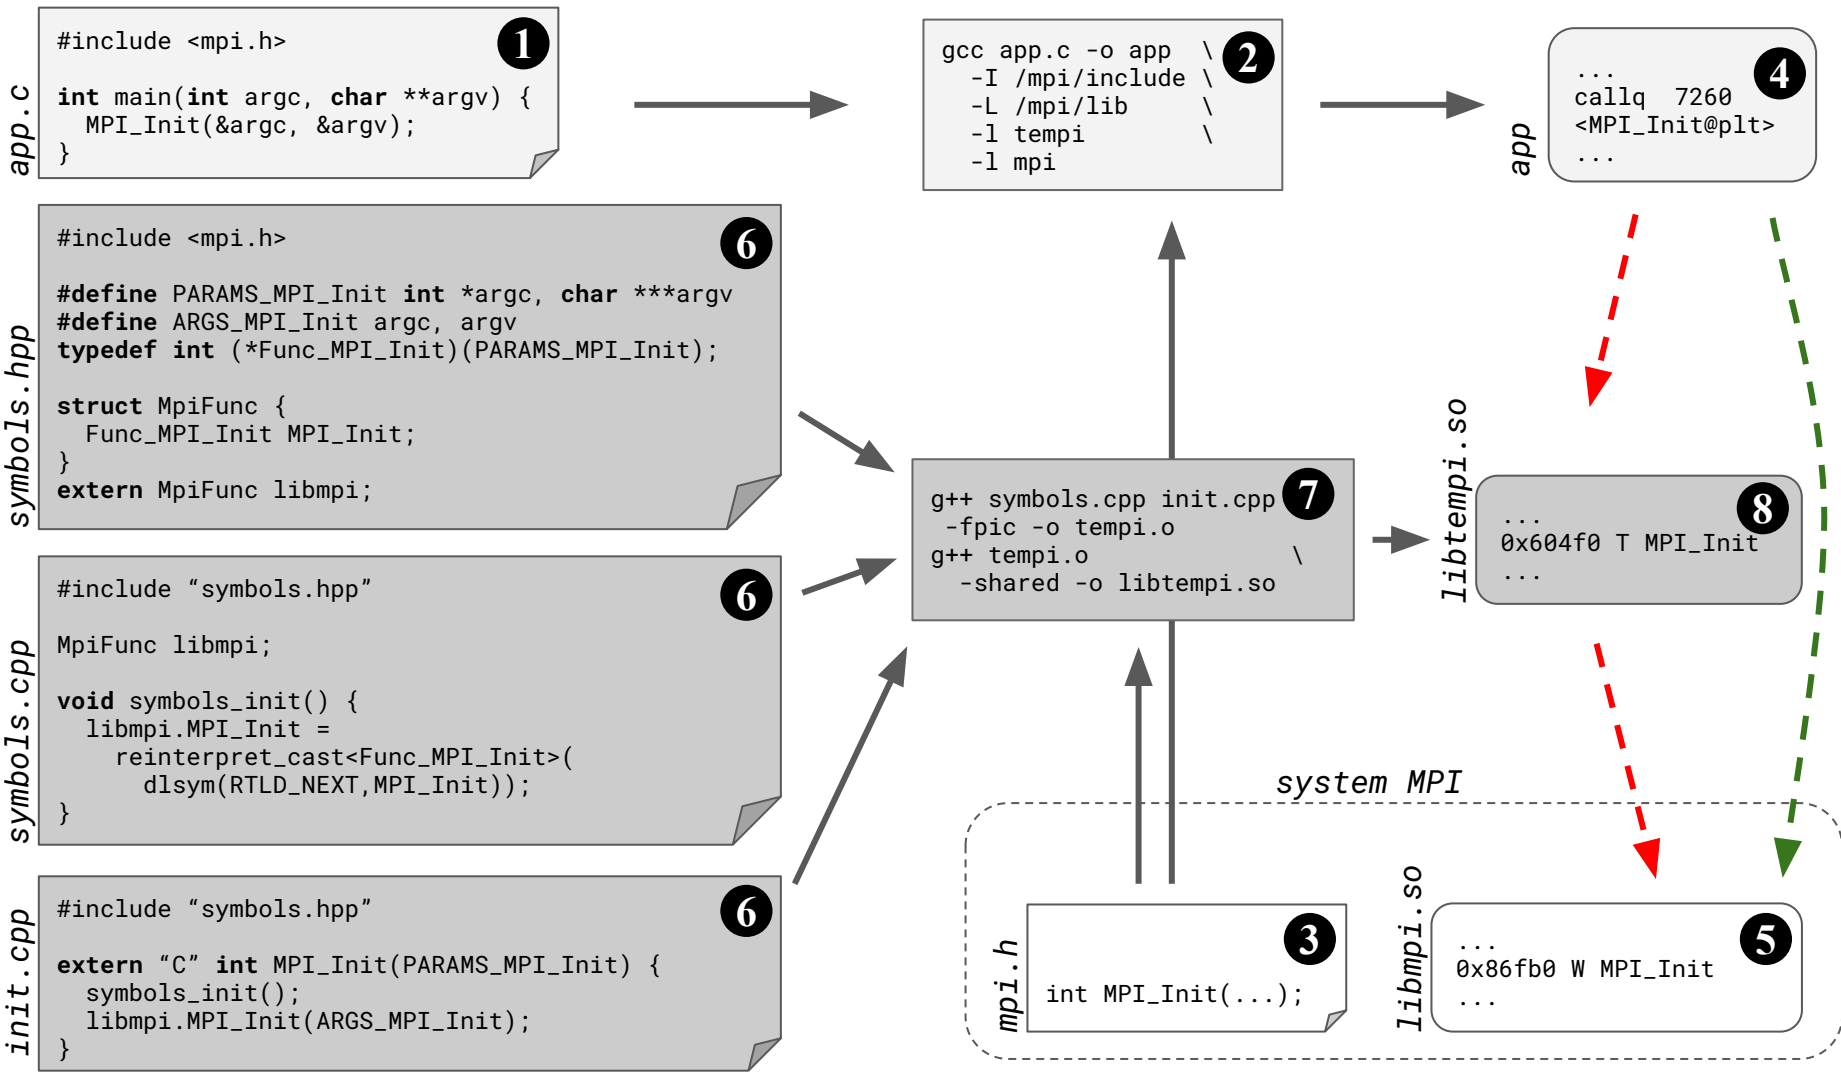
3

libmpi.so

```
...  
0x86fb0 W MPI_Init  
...
```

5





Conclusion

- Stencil code
- How they are parallelized
- Why non-contiguous data matters
- New MPI derived datatype approach
- A way to deploy experimental changes to systems.

Thank You

Slides & Paper: carlpearson.net > “Talks” > Click the May 5th link